

# Introduction

---

[NDC Melbourne Workshop](#) Event driven : All roads lead to Kafka

Build a taxi booking application with Apache Kafka.

During this workshop we'll discuss what Kafka is and how to use this exciting, modern distributed streaming platform in your next data project. From the ground up we'll build a practical data management system to stream data, aggregate results and deliver live visual mapping of a taxi booking application.

This hands-on labs you'll build an application that can

- Publish data to Kafka and subscribe to consume data from Kafka
- Serialize data in AVRO and Protobuf
- Register schemas and handle changes with schema evolution
- Leverage APIs in for producing and consuming messages against Kafka topics
- Build topologies to transform and aggregate streaming payloads
- Understand the Kafka connect framework to easily connect to external databases and elastic
- Build a live mapping dashboard with Kibana

This workshop is ideal for application developers, data engineers, or data scientists who need to interact with distributed messaging and streaming data platforms.

## Get started

---

### Prerequisites & setup

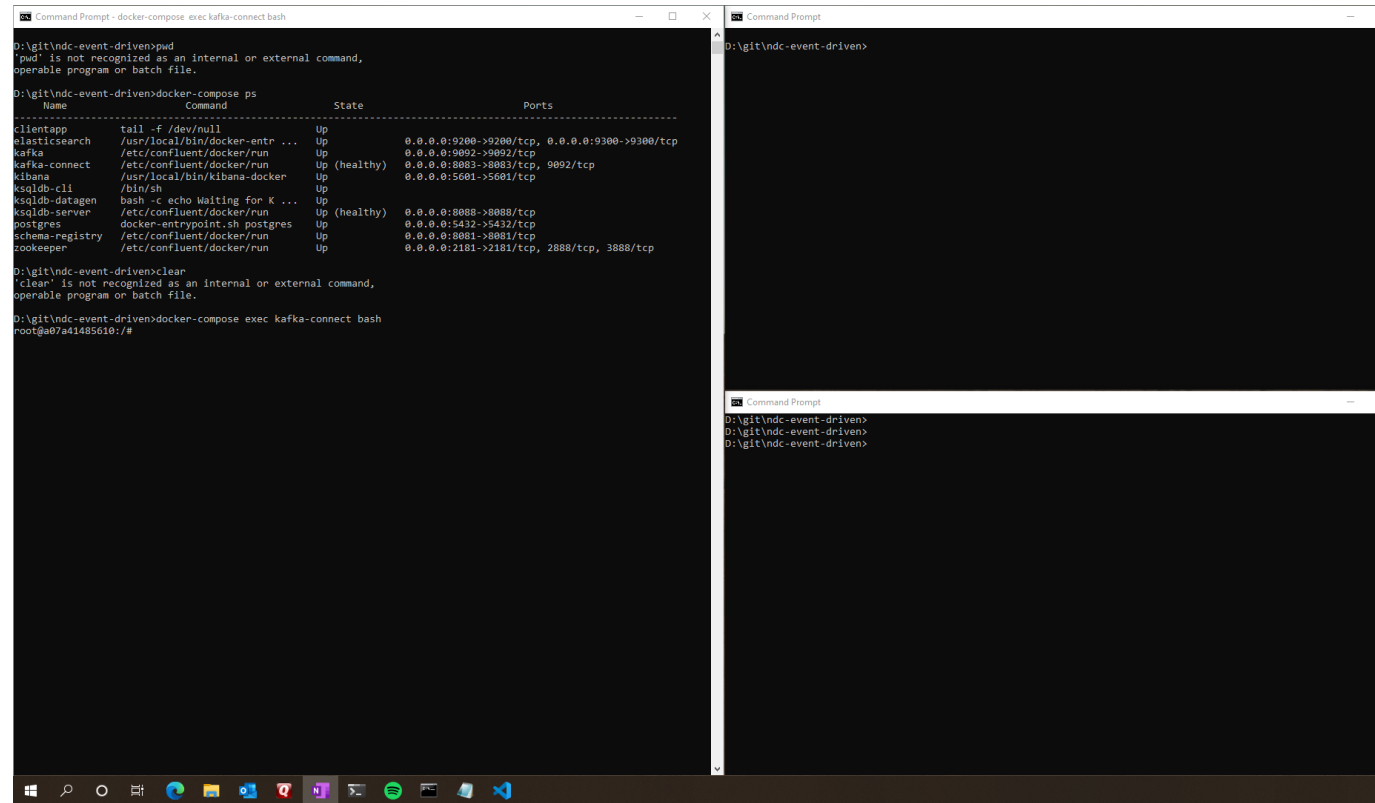
- install docker/docker-compose
- set your Docker maximum memory to something really big, such as 10GB. (preferences -> advanced -> memory)
- clone this repo!

```
mkdir ~/git
cd ~/git
git clone https://github.com/saubury/ndc-event-driven.git
cd ndc-event-driven
```

### How you'll work

You'll need *three* terminals for these exercises. Each one you should `cd ~/git/ndc-event-driven`

For simplicity arrange your three terminals so you can see the first and second at the same time (perhaps split horizontally if you're using iTerm2) Something like below.

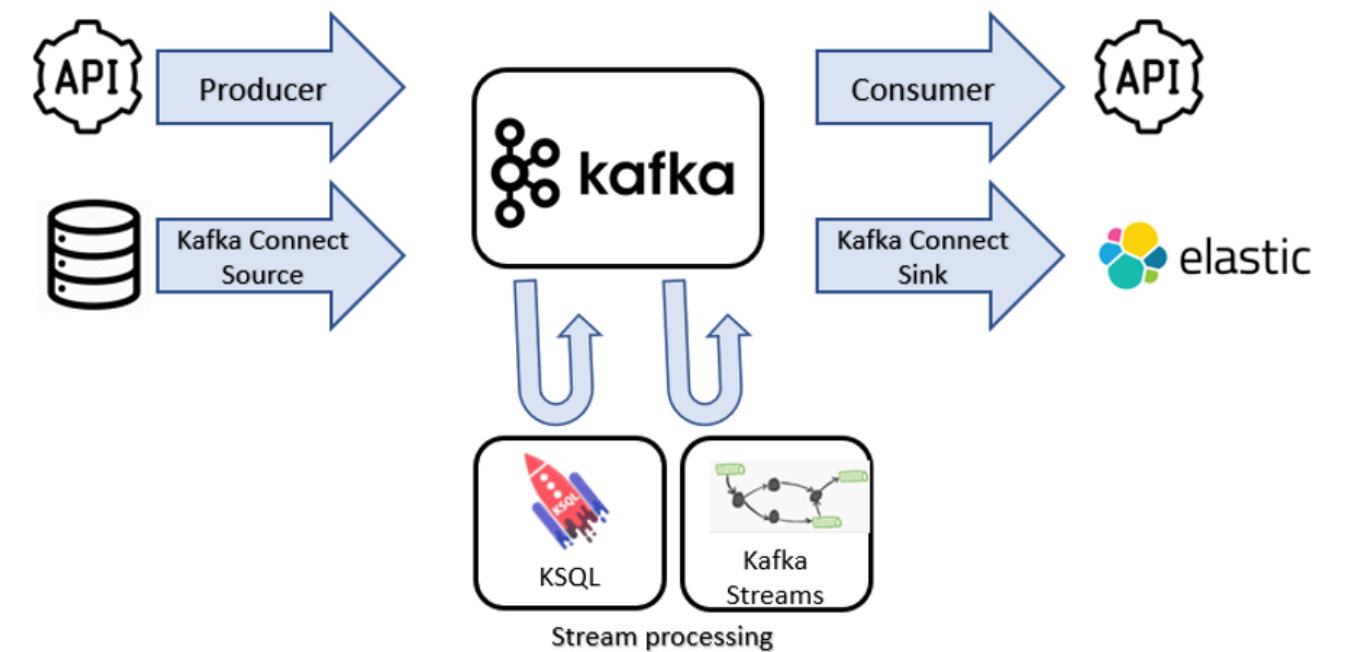


Docker Startup

Terminal 1

```
docker-compose up -d
```

Workshop 1 – Producers and Consumers



## Terminal 1

```
docker-compose exec kafka-connect bash
```

Create a topic

```
kafka-topics --bootstrap-server kafka:29092 --create --partitions 1 --replication-factor 1 --topic MYTOPIC
```

Check it's there

```
kafka-topics --list --bootstrap-server kafka:29092
```

Create a topic with 8 partitions

```
kafka-topics --bootstrap-server kafka:29092 --create --partitions 8 --replication-factor 1 --topic MYTOPIC8
```

Describe a topic - check it has 8 partitions

```
kafka-topics --describe --bootstrap-server kafka:29092 --topic MYTOPIC8
```

Create a topic with replication greater than number of brokers

```
kafka-topics --bootstrap-server kafka:29092 --create --replication-factor 2 --topic MYTOPIC_REPLICATED
```

*Create a Producer*

Write some text from STDIN

```
kafka-console-producer --broker-list kafka:29092 --topic MYTOPIC
```

Now type some things into first (original) terminal (and press ENTER).

## Terminal 2

*Create a Consumer*

Start another terminal

```
docker-compose exec kafka-connect bash
```


In the newly created (second) terminal let's start reading from new Kafka topic

```
kafka-console-consumer --bootstrap-server kafka:29092 --topic MYTOPIC --from-  
beginning
```

Each line you type in the first terminal should appear in second terminal

What have we learnt? It's easy to be a producer or consumer. Out of the box Kafka doesn't care what you're writing - it's just a bunch of bytes

## Structured Data with AVRO



### Terminal 1

At UNIX prompt, (Note: Press Ctrl C to exit the producer script)

```
kafka-topics --bootstrap-server kafka:29092 --create --partitions 1 --replication-  
factor 1 --topic COMPLAINTS_AVRO
```

```
kafka-avro-console-producer --broker-list kafka:29092 --property  
schema.registry.url="http://schema-registry:8081" --topic COMPLAINTS_AVRO \  
--property value.schema='  
{  
  "type": "record",  
  "name": "myrecord",  
  "fields": [  
    {"name": "customer_name", "type": "string" }  
    , {"name": "complaint_type", "type": "string" }  
    , {"name": "trip_cost", "type": "float" }  
    , {"name": "new_customer", "type": "boolean"}  
  ]  
}
```

```
]
}' << EOF
{"customer_name":"Carol", "complaint_type":"Late arrival", "trip_cost": 19.60,
 "new_customer": false}
EOF
```

## Terminal 2

Press Ctrl+C and Ctrl+D and run the following curl command.

BTW, this is AVRO

```
curl -s -X GET http://localhost:8081/subjects/COMPLAINTS_AVRO-value/versions/1
```

## AVRO Schema Evolution

Let's add a loyalty concept to our complaints topic - we'll add "number\_of\_rides" to the payload

## Terminal 1

```
kafka-avro-console-producer --broker-list kafka:29092 --property
schema.registry.url="http://schema-registry:8081" --topic COMPLAINTS_AVRO \
--property value.schema='
{
  "type": "record",
  "name": "myrecord",
  "fields": [
    {"name": "customer_name", "type": "string" }
    , {"name": "complaint_type", "type": "string" }
    , {"name": "trip_cost", "type": "float" }
    , {"name": "new_customer", "type": "boolean"}
    , {"name": "number_of_rides", "type": "int", "default" : 1}
  ]
}' << EOF
{"customer_name":"Ed", "complaint_type":"Dirty car", "trip_cost": 29.10,
 "new_customer": false, "number_of_rides": 22}
EOF
```

## Terminal 2

Let's see what schemas we have registered now

```
curl -s -X GET http://localhost:8081/subjects/COMPLAINTS_AVRO-value/versions

curl -s -X GET http://localhost:8081/subjects/COMPLAINTS_AVRO-value/versions/1 |
jq '.'
```

```
curl -s -X GET http://localhost:8081/subjects/COMPLAINTS_AVRO-value/versions/2 |  
jq '.'
```

or you can also use:

```
curl -s -X GET http://localhost:8081/subjects/COMPLAINTS_AVRO-value/versions/2 |  
jq -r .schema | jq .
```

## Workshop 2 – Kafka Connect

---

Let's copy data from an upstream database which has a list of ride users. Connecting Kafka to and from other systems (such as a database or object store) is a very common task. The Kafka Connect framework has a plug in archecture which allows you to *source* from an upstream system or *sink* into a downstream system.

### Setup Postgres source database

#### in Terminal 1:

Exit the kafka-connect container by pressing Ctrl+D.

```
cat scripts/postgres-setup.sql  
  
docker-compose exec postgres psql -U postgres -f /scripts/postgres-setup.sql
```

To look at the Postgres table

```
docker-compose exec postgres psql -U postgres -c "select * from users;"
```

### Kafka Connect Setup

Our goal now is to source data continuously from our Postgres database and produce into Kafka. We'll use Kafka connect as the framework, and a JDBC Postgres Source connector coto connect to the database

Have a look at [scripts/connect\\_source\\_postgres.json](#)

Load connect config

```
curl -k -s -S -X PUT -H "Accept: application/json" -H "Content-Type:  
application/json" --data @./scripts/connect_source_postgres.json  
http://localhost:8083/connectors/src_pg/config
```

```
curl -s -X GET http://localhost:8083/connectors/src_pg/status | jq '.'
```

## Terminal 2

Now let's consume the topic by starting a consumer inside the kafka-connect container to

```
kafka-avro-console-consumer --bootstrap-server kafka:29092 --topic db-users --
from-beginning --property schema.registry.url="http://schema-registry:8081"
```

## Terminal 1

Insert a new database row into Postgres

```
docker exec -it postgres psql -U postgres -c "INSERT INTO users (userid, username)
VALUES ('J', 'Jane');" 
```

You *should* see Jane arrive automatically into the Kafka topic

# Workshop 3 – Stream Processing

---

## Generate ride request data

Create a stream of rider requests

## Terminal 3

```
docker-compose exec ksqldb-datagen ksql-datagen schema=/scripts/riderequest.avro
format=avro topic=riderequest key=rideid msgRate=1 iterations=10 bootstrap-
server=kafka:29092 schemaRegistryUrl=http://schema-registry:8081 value-format=avro
```

In Terminal 2, (Exit the existing consumer by pressing Ctrl+C) Check the AVRO output of the **riderequest** topic. Press ^C when you've seen a few records.

```
kafka-avro-console-consumer --bootstrap-server kafka:29092 --topic riderequest --
from-beginning --property schema.registry.url="http://schema-registry:8081"
```

## Build a stream processor

We have a constant stream of rider requests arriving in the **riderequest** topic. But each request has only a **userid** (such as **J**) and no name (like **Jane**). Also, the rider location has separate latitude and longitude fields; we want to be able to join them together as single string field (to form a geom - **cast(rr.LATITUDE as varchar) || ', ' || cast(rr.LONGITUDE as varchar)**)

Let's build a stream processor to consume from the `riderequest` topic and `db-users` topics, join them and produce into a new topic along with a new location attribute.

Will build our stream processor in ksql.

## ksqlDB CLI

Build a stream processor

### Terminal 2

```
docker-compose exec ksqldb-cli ksql http://ksqldb-server:8088
```

Run the KSQL script:

```
ksql
ksql> run script '/scripts/join_topics.ksql';
exit;
```

And if you want to check

**Terminal 1** from inside the Kafka-connect container

```
kafka-console-consumer --bootstrap-server kafka:29092 --topic RIDESANDUSERSJSON
```

## Workshop 4 – Visualizations

---

### Sink to Elastic/Kibana

Setup dynamic elastic templates

### Terminal 2

At the console prompt

```
./scripts/load_elastic_dynamic_template
```

Now we need a sink connector to read from the topic RIDESANDUSERSJSON

Load connect config.



```
curl -k -s -S -X PUT -H "Accept: application/json" -H "Content-Type: application/json" --data @./scripts/connect_sink_elastic.json http://localhost:8083/connectors/sink_elastic/config
```

```
curl -s -X GET http://localhost:8083/connectors/sink_elastic/status | jq '.'
```

## Kibana Dashboard Import

- Navigate to <http://localhost:5601/app/kibana#/management/kibana/objects>
- Click Import
- Select file 06\_kibana\_export.json
- Click Automatically overwrite all saved objects? and select Yes, overwrite all objects
- Kibana - Open Dashboard
- Open <http://localhost:5601/app/kibana#/dashboards>

## Workshop 5 – Producers & Protobuf

---

Build a Python producer

### Application Binding - Protobuf classes with Python

Let us now build an application demonstrating protobuf classes.

```
docker-compose exec clientapp bash
```

### Generate protobuf classes

To generate protobuf classes you must first install the protobuf compiler [protoc](#). See the [protocol buffer docs](#) for instructions on installing and using protoc.

#### Python compile schema

```
cd /python-app  
cat ./meal.proto
```

```
protoc -I=. --python_out=. ./meal.proto
```

This will create the `meal_pb2.py` Python class file.

You can now build protobuf classes and produce into Kafka with code like this

```
pip install -r requirements.txt
```

Review the producer code. Note the `meal_pb2.py` Python class file is included. This class is used to construct a payload in protobuf. The producer will send the schema to the schema registry and the payload to the `MEAL_DELIVERY` topic.

```
cat producer-protobuf.py
```

Run the producer

```
python producer-protobuf.py
```

## Review the topic

Now to check the contents of the `MEAL_DELIVERY` topic. We will use the `kafka-protobuf-console-consumer` command line consumer tool to inspect the protobuf payload.

```
docker-compose exec kafka-connect bash
```

```
kafka-protobuf-console-consumer --bootstrap-server kafka:29092 --topic  
MEAL_DELIVERY --from-beginning --property schema.registry.url="http://schema-  
registry:8081"
```

## Verify schema

Check this is Protobuf

```
curl -s -X GET http://localhost:8081/subjects/MEAL_DELIVERY-value/versions/1
```